

Scaling Data Analysis

Google Cloud Platform Fundamentals: Big Data and Machine Learning

Version #1.1



Agenda



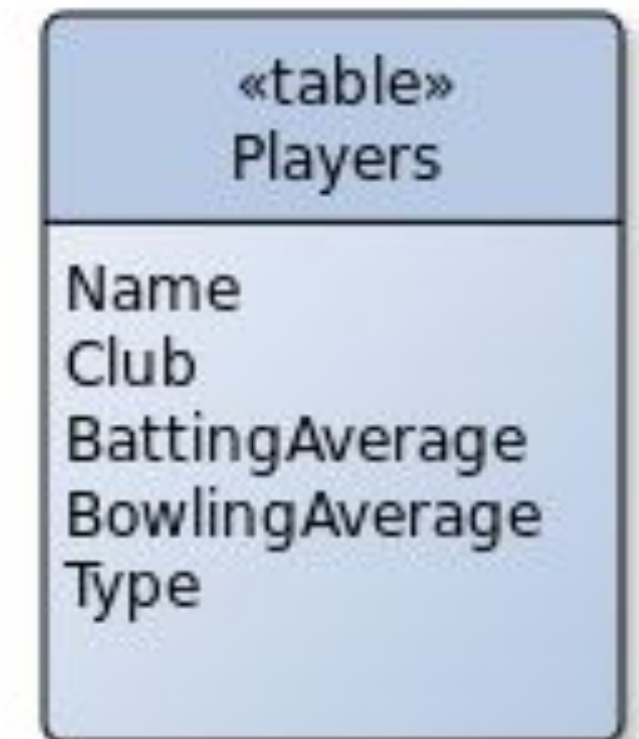
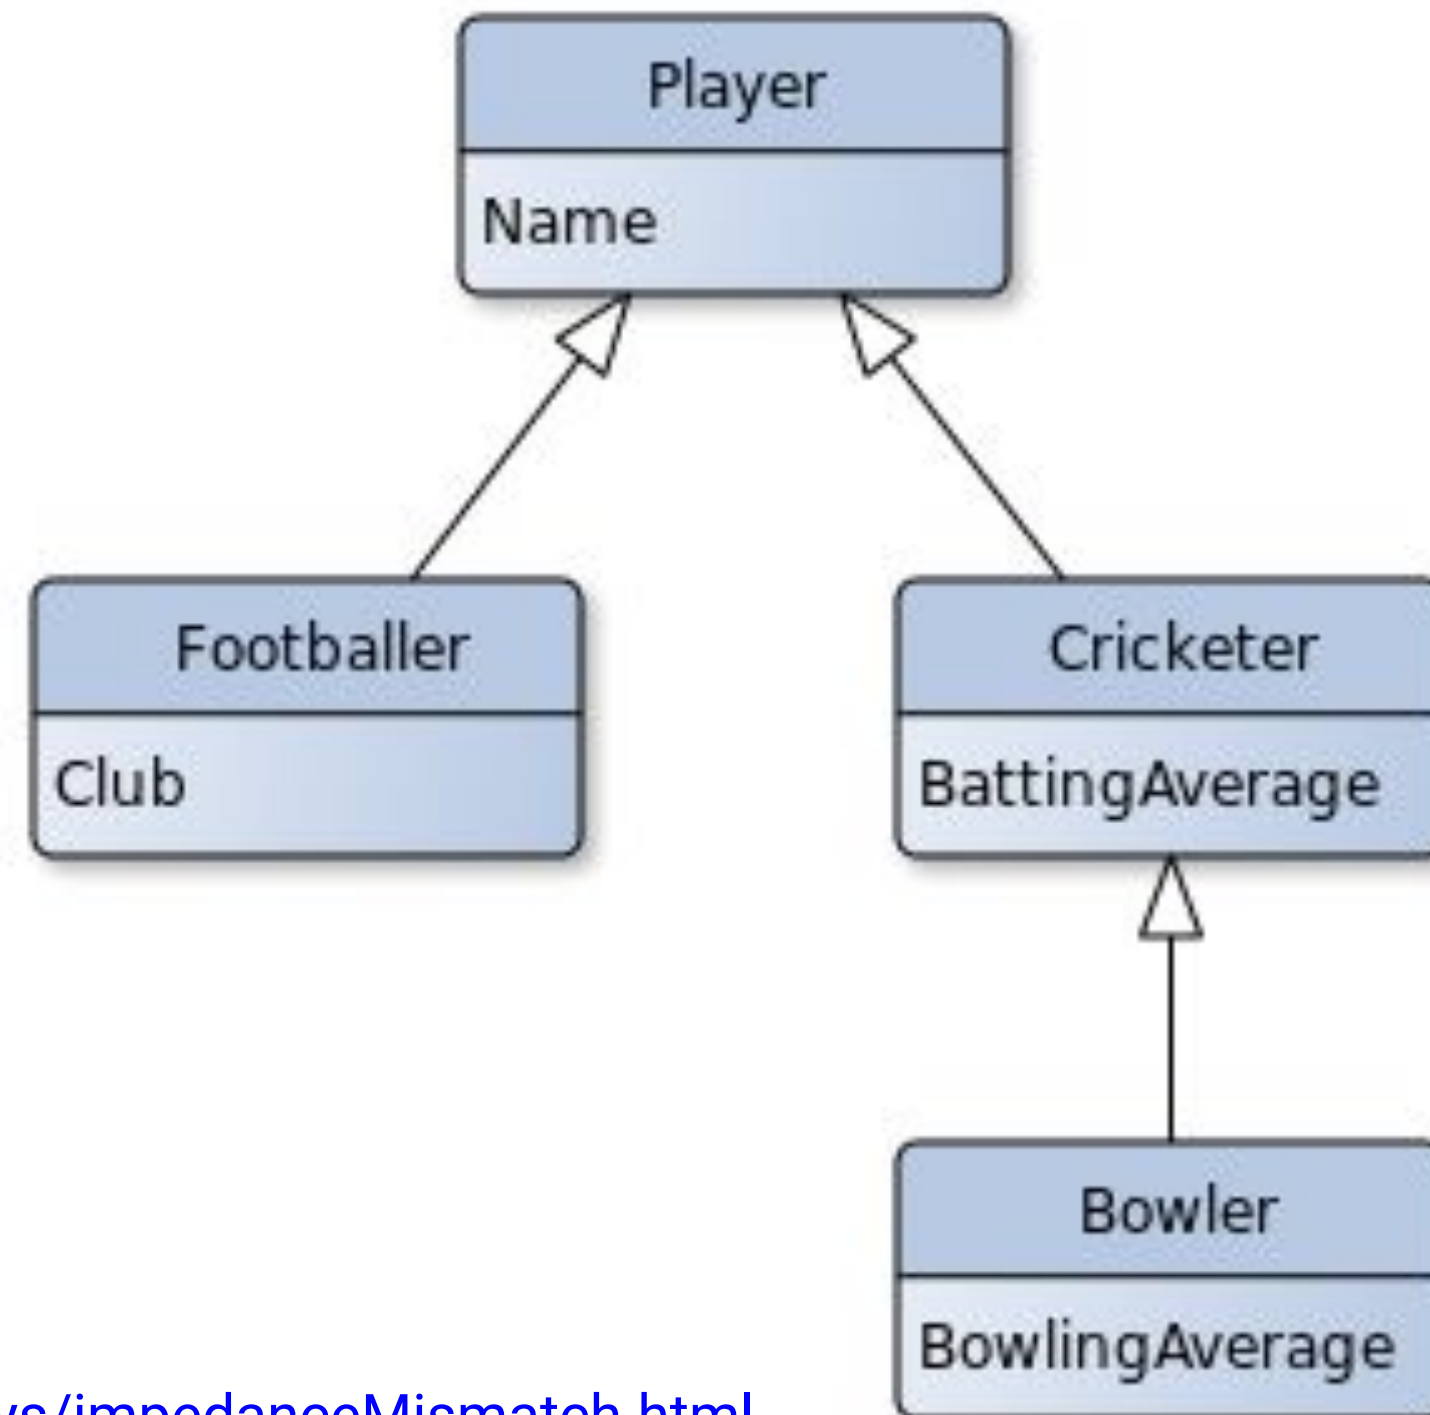
Agenda

Fast random access

Warehouse and interactively query petabytes

Interactive, iterative development + Lab

Relational tables are hard to use from object-oriented programs



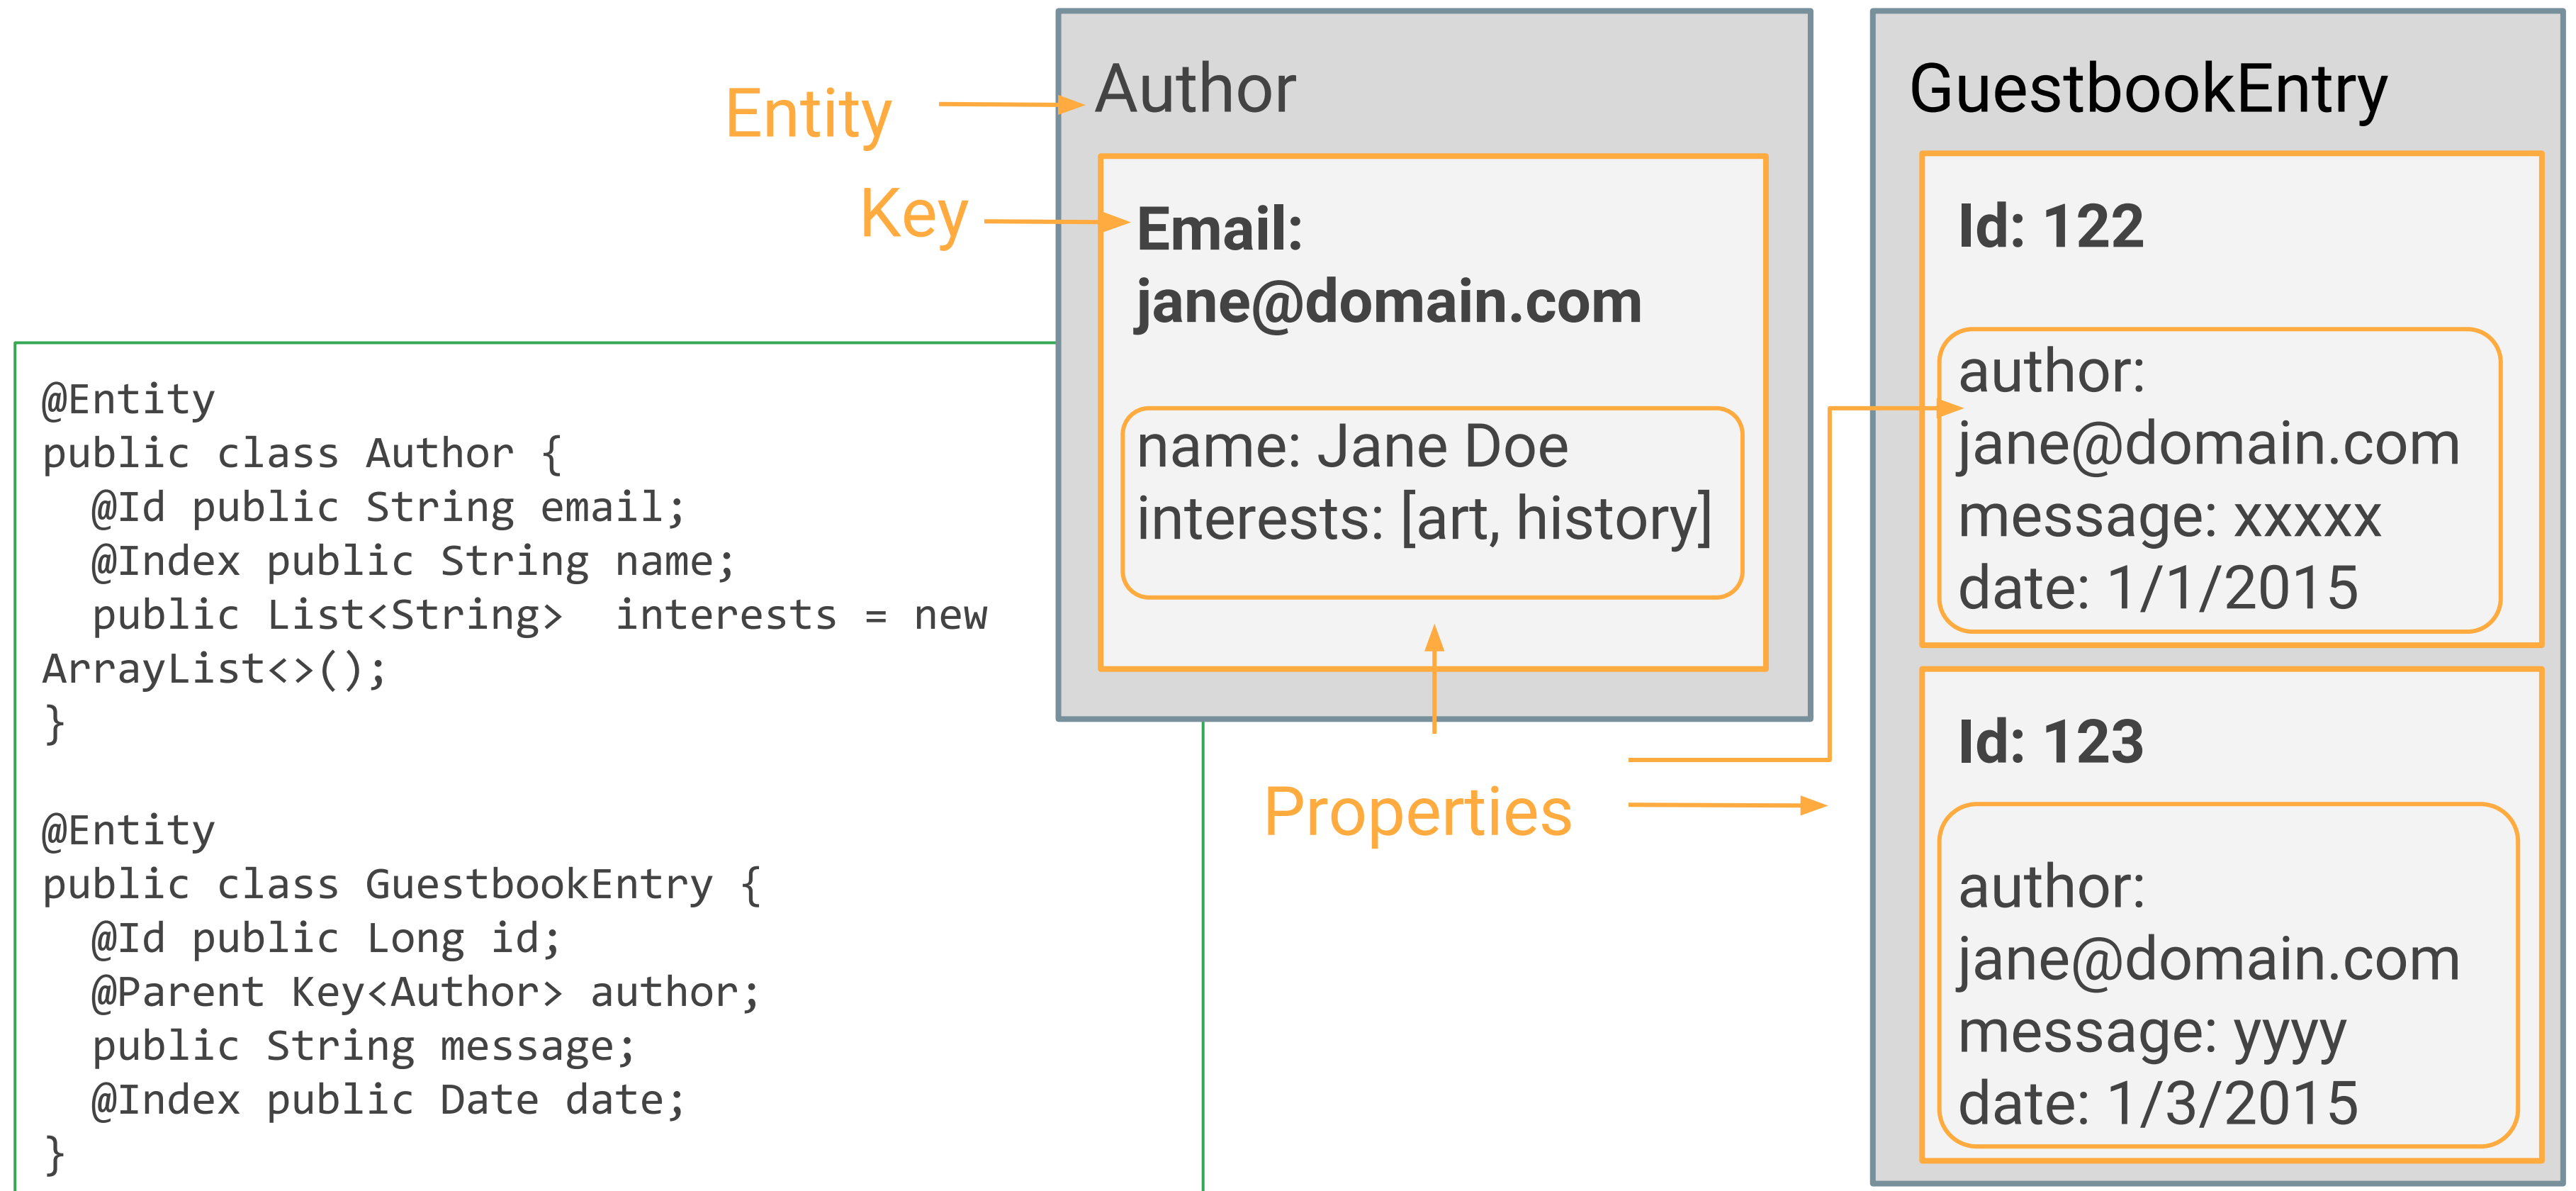
<http://www.agiledata.org/essays/impedanceMismatch.html>

https://en.wikipedia.org/wiki/Object_database

Choose storage based on access pattern

	Cloud Storage	Cloud SQL	Datastore	Bigtable	BigQuery
Capacity	Petabytes +	Gigabytes	Terabytes	Petabytes	Petabytes
Access metaphor	Like files in a file system	Relational database	Persistent Hashmap	Key-value(s), HBase API	Relational
Read	Have to copy to local disk	SELECT rows	filter objects on property	scan rows	SELECT rows
Write	One file	INSERT row	put object	put row	Batch/stream
Update granularity	An object (a "file")	Field	Attribute	Row	Field
Usage	Store blobs	No-ops SQL database on the cloud	Structured data from AppEngine apps	No-ops, high throughput, scalable, flattened data	Interactive SQL* querying fully managed warehouse

Datastore is like a persistent HashMap



CRUD operations are easily implemented on Datastore

```
@Entity
public class Author {
    @Id public String email;
    @Index public String name;
    public List<String> interests =
new ArrayList<>();
}
```

// CREATE

```
Author xjin = new Author("xjin@bu.edu", "Ha Jin");
xjin.interests.add("Misty Poetry");
ofy.save().entity(xjin);
```

// READ

```
Iterable<Author> authors =
ofy().load().type(Author.class).filter("name", "Ha Jin");

Author jh = ofy.load().type(Author.class).id("xjin@bu.edu").now();
```

// UPDATE

```
jh.name = "Jīn Xuěfēi (金雪飞)";
ofy().save().entity(jh).now();
```

// DELETE

```
ofy().delete().entity(jh).now();
```

Choose storage product based on access pattern

	Cloud Storage	Cloud SQL	Datastore	Bigtable	BigQuery
Capacity	Petabytes +	Gigabytes	Terabytes	Petabytes	Petabytes
Access metaphor	Like files in a file system	Relational database	Persistent Hashmap	Key-value(s), HBase API	Relational
Read	Have to copy to local disk	SELECT rows	filter objects on property	scan rows	SELECT rows
Write	One file	INSERT row	put object	put row	Batch/stream
Update granularity	An object (a "file")	Field	Attribute	Row (write new row instead)	Field
Usage	Store blobs	No-ops SQL database on the cloud	Structured data from AppEngine apps	No-ops, high throughput, scalable, flattened data	Interactive SQL* querying fully managed warehouse

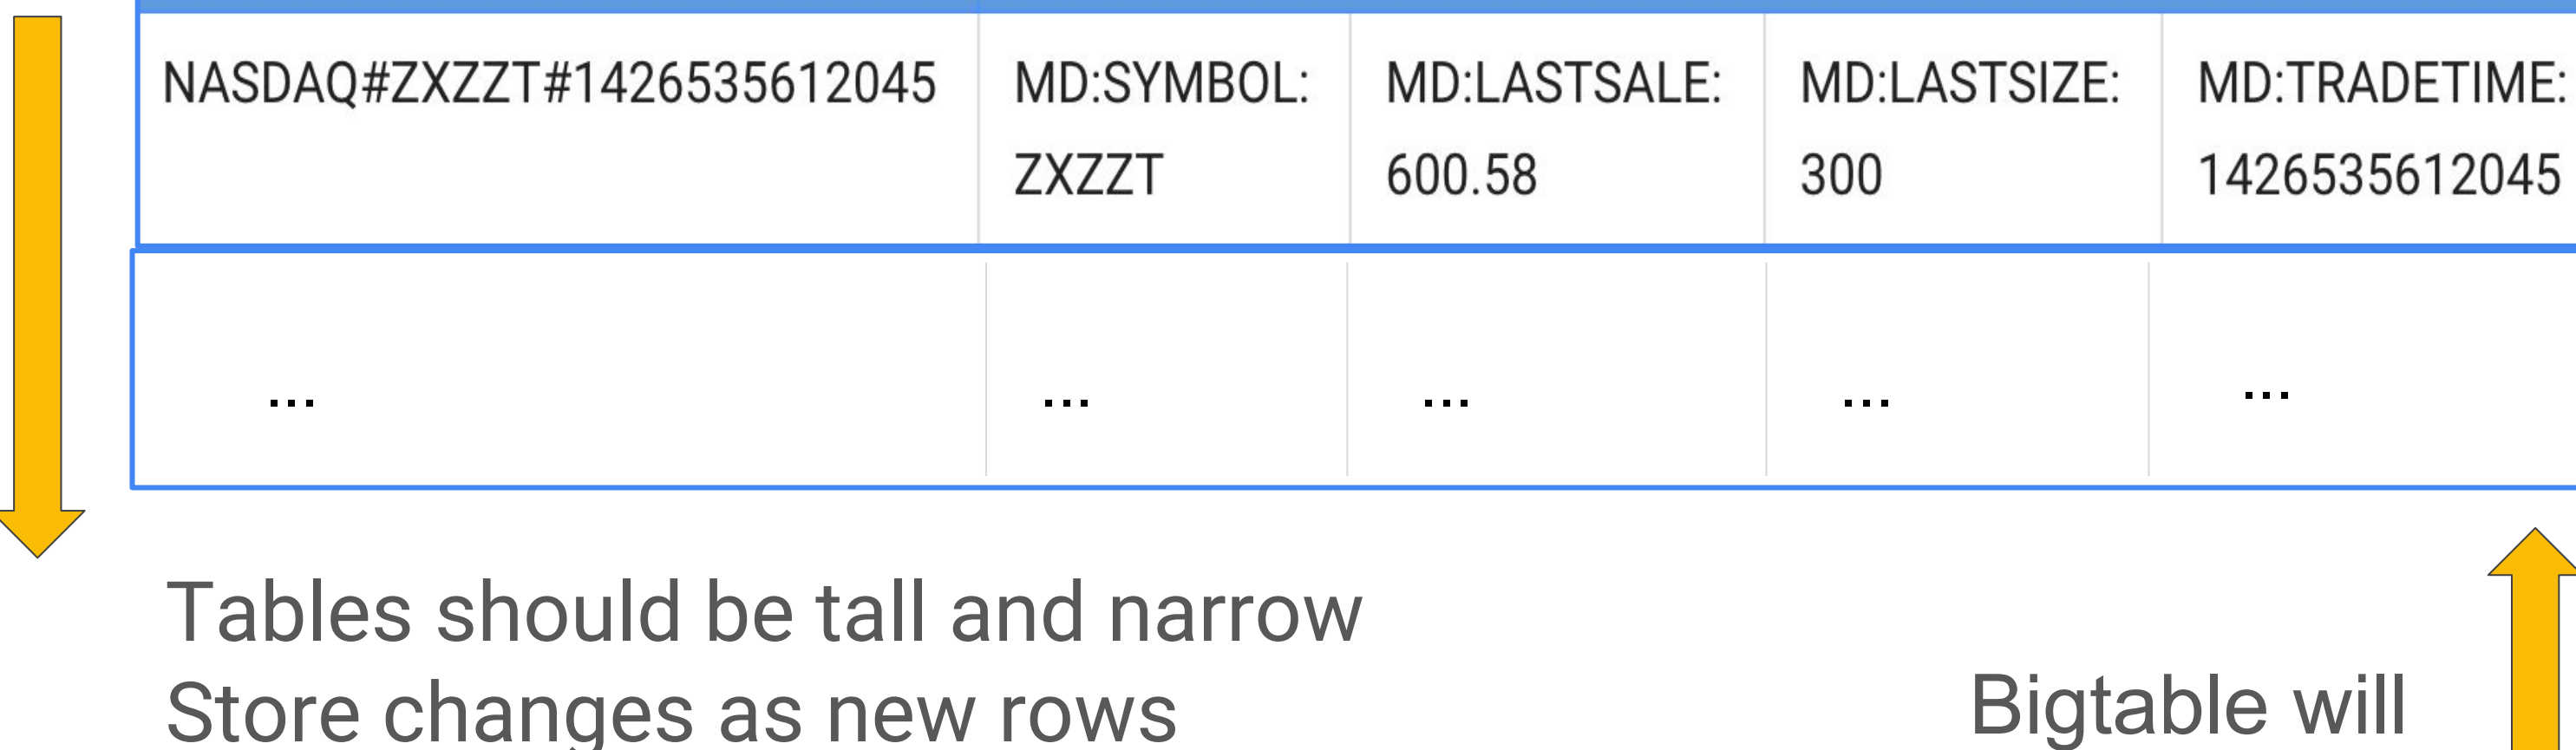
Bigtable is meant for high throughput data where access is primarily for a range of Row Key prefixes

Design row key with most common query in mind

Row key	Column data				
NASDAQ#ZXZZT#1426535612045	MD:SYMBOL: ZXZZT	MD:LASTSALE: 600.58	MD:LASTSIZE: 300	MD:TRADETIME: 1426535612045	MD:EXCHANGE: NASDAQ

Design row key to minimize hotspots

Bigtable is meant for high throughput data where access is primarily for a range of Row Key prefixes



Row key	Column data				
NASDAQ#ZXZZT#1426535612045	MD:SYMBOL: ZXZZT	MD:LASTSALE: 600.58	MD:LASTSIZE: 300	MD:TRADETIME: 1426535612045	MD:EXCHANGE: NASDAQ
...

Tables should be tall and narrow
Store changes as new rows

Bigtable will
automatically compact
the table

Short meaningful column names reduce storage and RPC overhead

Design row key with most common query in mind

Column families is a quick way to get some hierarchy

Row key	Column data				
NASDAQ#ZXZZT#1426535612045	MD:SYMBOL: ZXZZT	MD:LASTSALE: 600.58	MD:LASTSIZE: 300	MD:TRADETIME: 1426535612045	MD:EXCHANGE: NASDAQ

Design row key to minimize hotspots

Use short column names
Designed for sparse tables

Can work with Bigtable using the HBase API

```
import org.apache.hadoop.hbase.*;
import org.apache.hadoop.hbase.client.*;
import org.apache.hadoop.hbase.util.*;

byte[] CF = Bytes.toBytes("MD"); // column family
Connection connection = ConnectionFactory.createConnection(...)
Table table = null;
try {
    table = connection.getTable(TABLE_NAME);
    Put p = new Put(Bytes.toBytes("NASDAQ#GOOG #1234561234561"));
    p.addColumn(CF, Bytes.toBytes("SYMBOL"), Bytes.toBytes("GOOG"));

    p.addColumn(CF, Bytes.toBytes("LASTSALE"), Bytes.toBytes("742.03d"));

    ...
    table.put(p);
} finally {
    if (table != null) table.close();
}
```

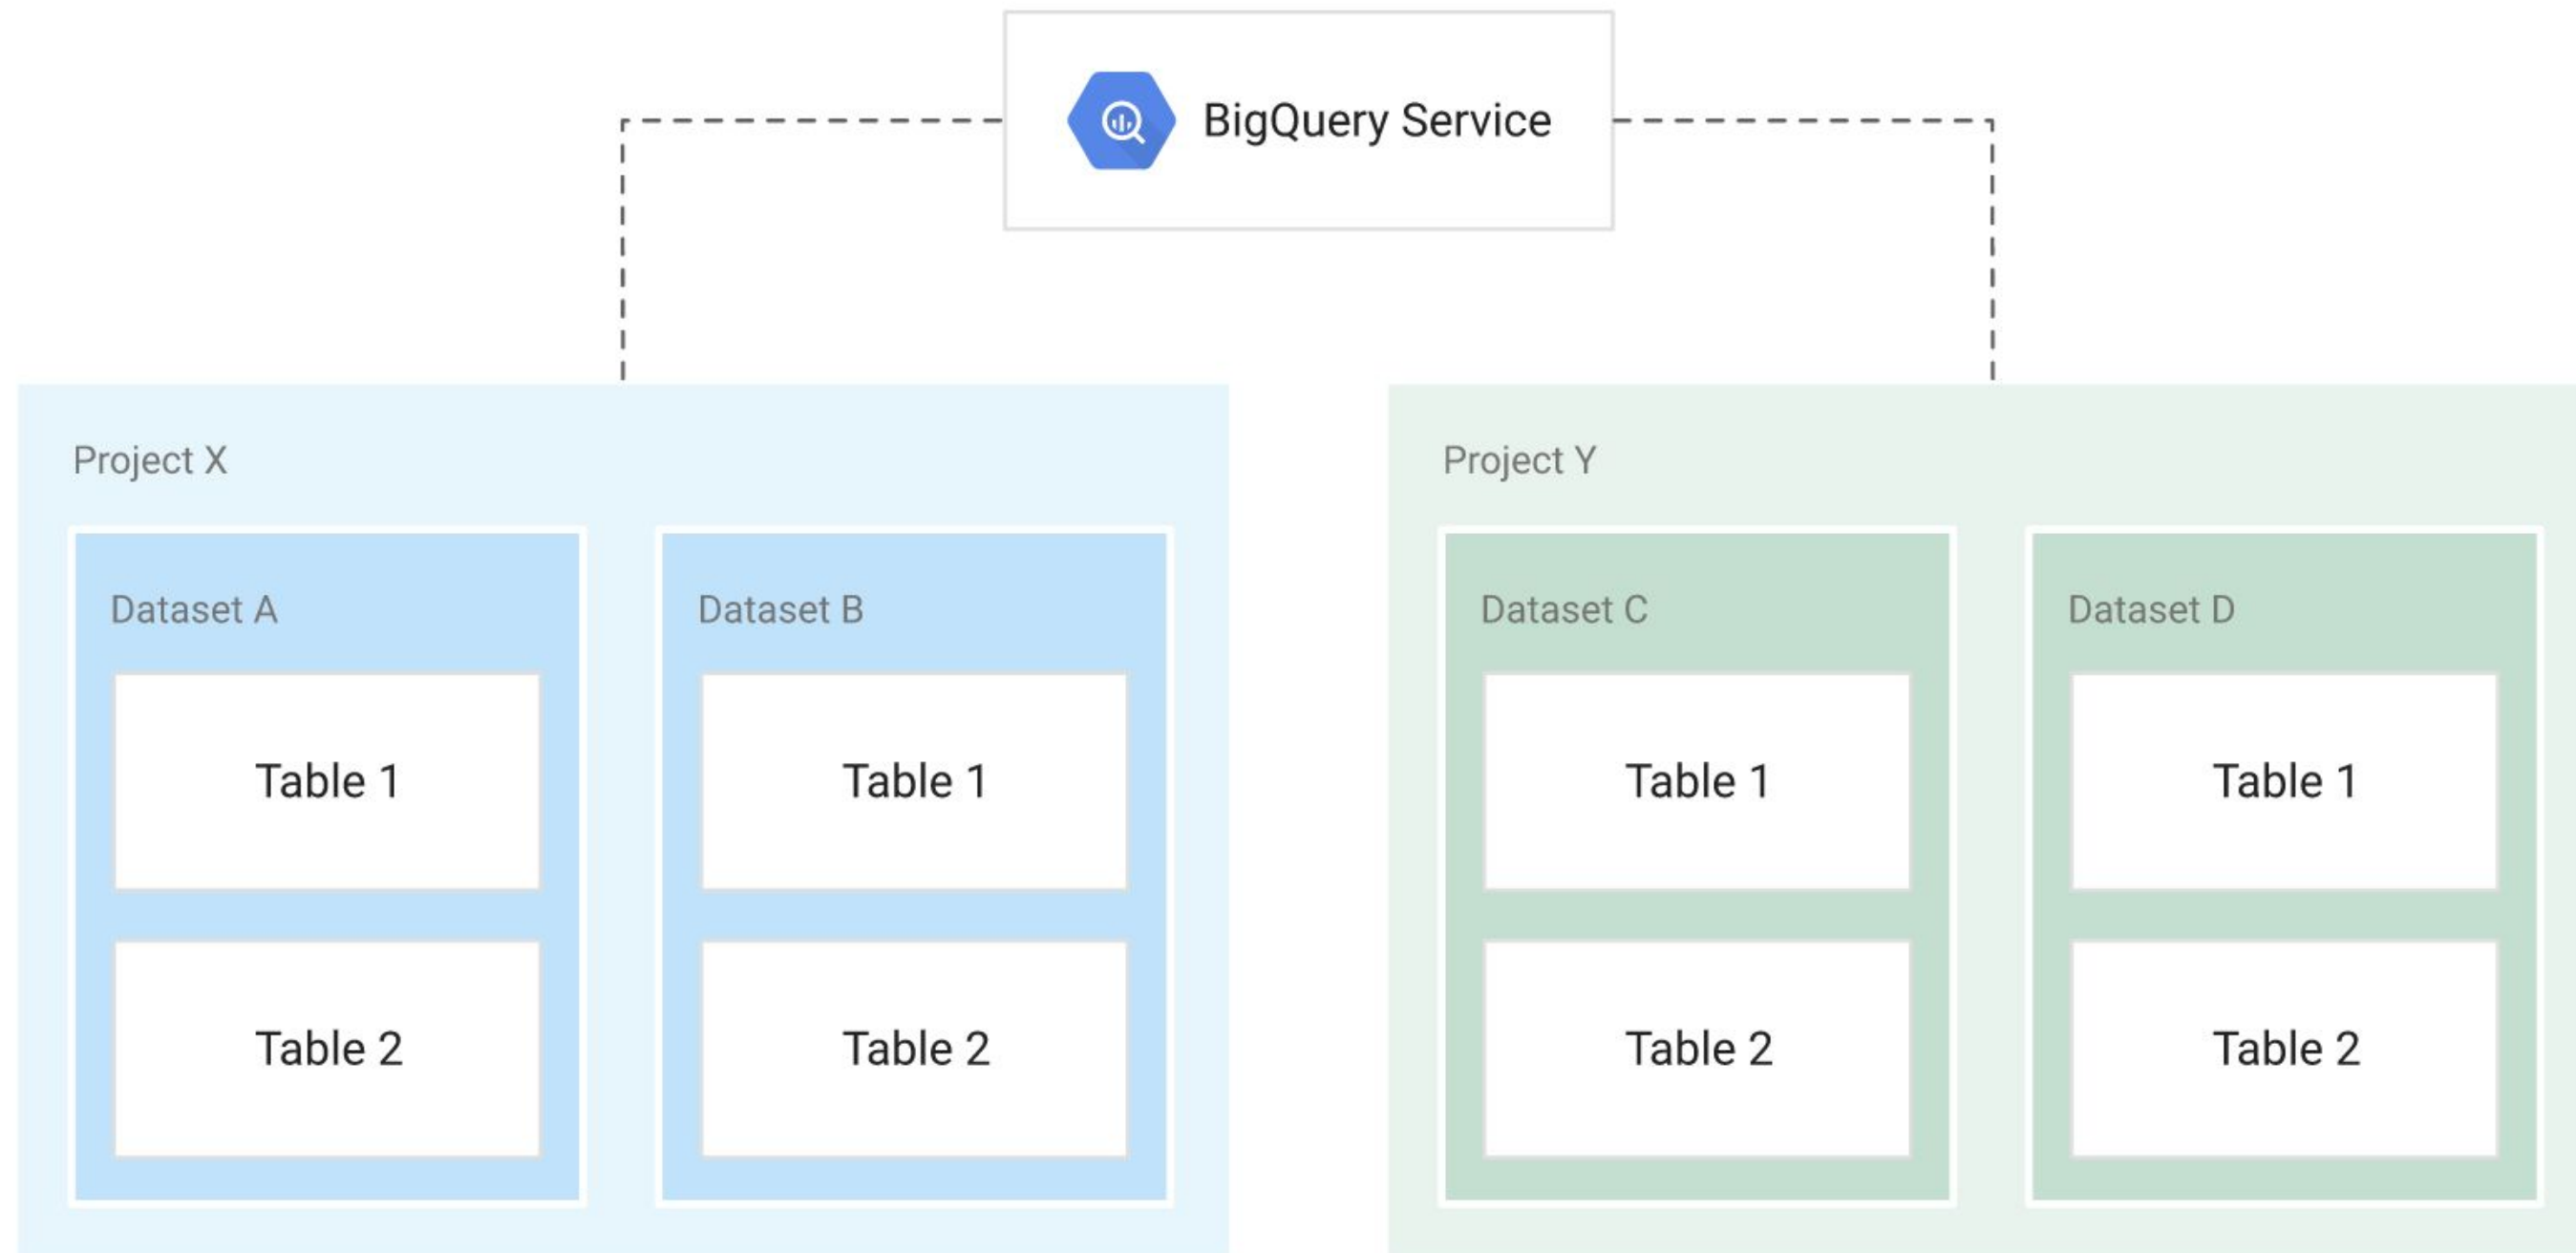
Agenda

Fast random access

Warehouse and interactively query petabytes

Interactive, iterative development + Lab

BigQuery is a fully managed data warehouse that lets you do ad-hoc SQL queries on massive volumes of data



A demo of BigQuery on a 10 billion-row dataset shows what it is and what it can do

```
#standardsql
SELECT
  language, SUM(views) as views
FROM `bigquery-samples.wikipedia_benchmark.Wiki10B`
WHERE
  title like "%google%"
GROUP by language
ORDER by views DESC
```


Familiar, SQL 2011 query language

Interactive ad-hoc analysis of
petabyte-scale databases


No need to provision clusters

Three ways of loading data into BigQuery


Files on disk or Cloud Storage



Stream Data



Federated data source



With Federated data sources, you can directly query files on Cloud Storage, without having to ingest them into BigQuery

Create Table

Source Data ☒ Create from source ☐ Create empty table

Repeat job

Select Previous Job

?

Location

Google Cloud Storage

gs://cloud-training-demos/flights/raw/201601.csv

?

File format

CSV

[View Files](#)

Also: JSON/Avro/Google Sheet

Destination Table

Table name

flights

jan2016

?

Table type

External table

?

Schema ☒ Automatically detect

?

Schema will be automatically generated.

Can also pass in a schema

Options

Header rows to skip

0

?

Number of errors allowed

0

?

Create Table

Agenda

Fast random access

Warehouse and interactively query petabytes

Interactive, iterative development + Lab

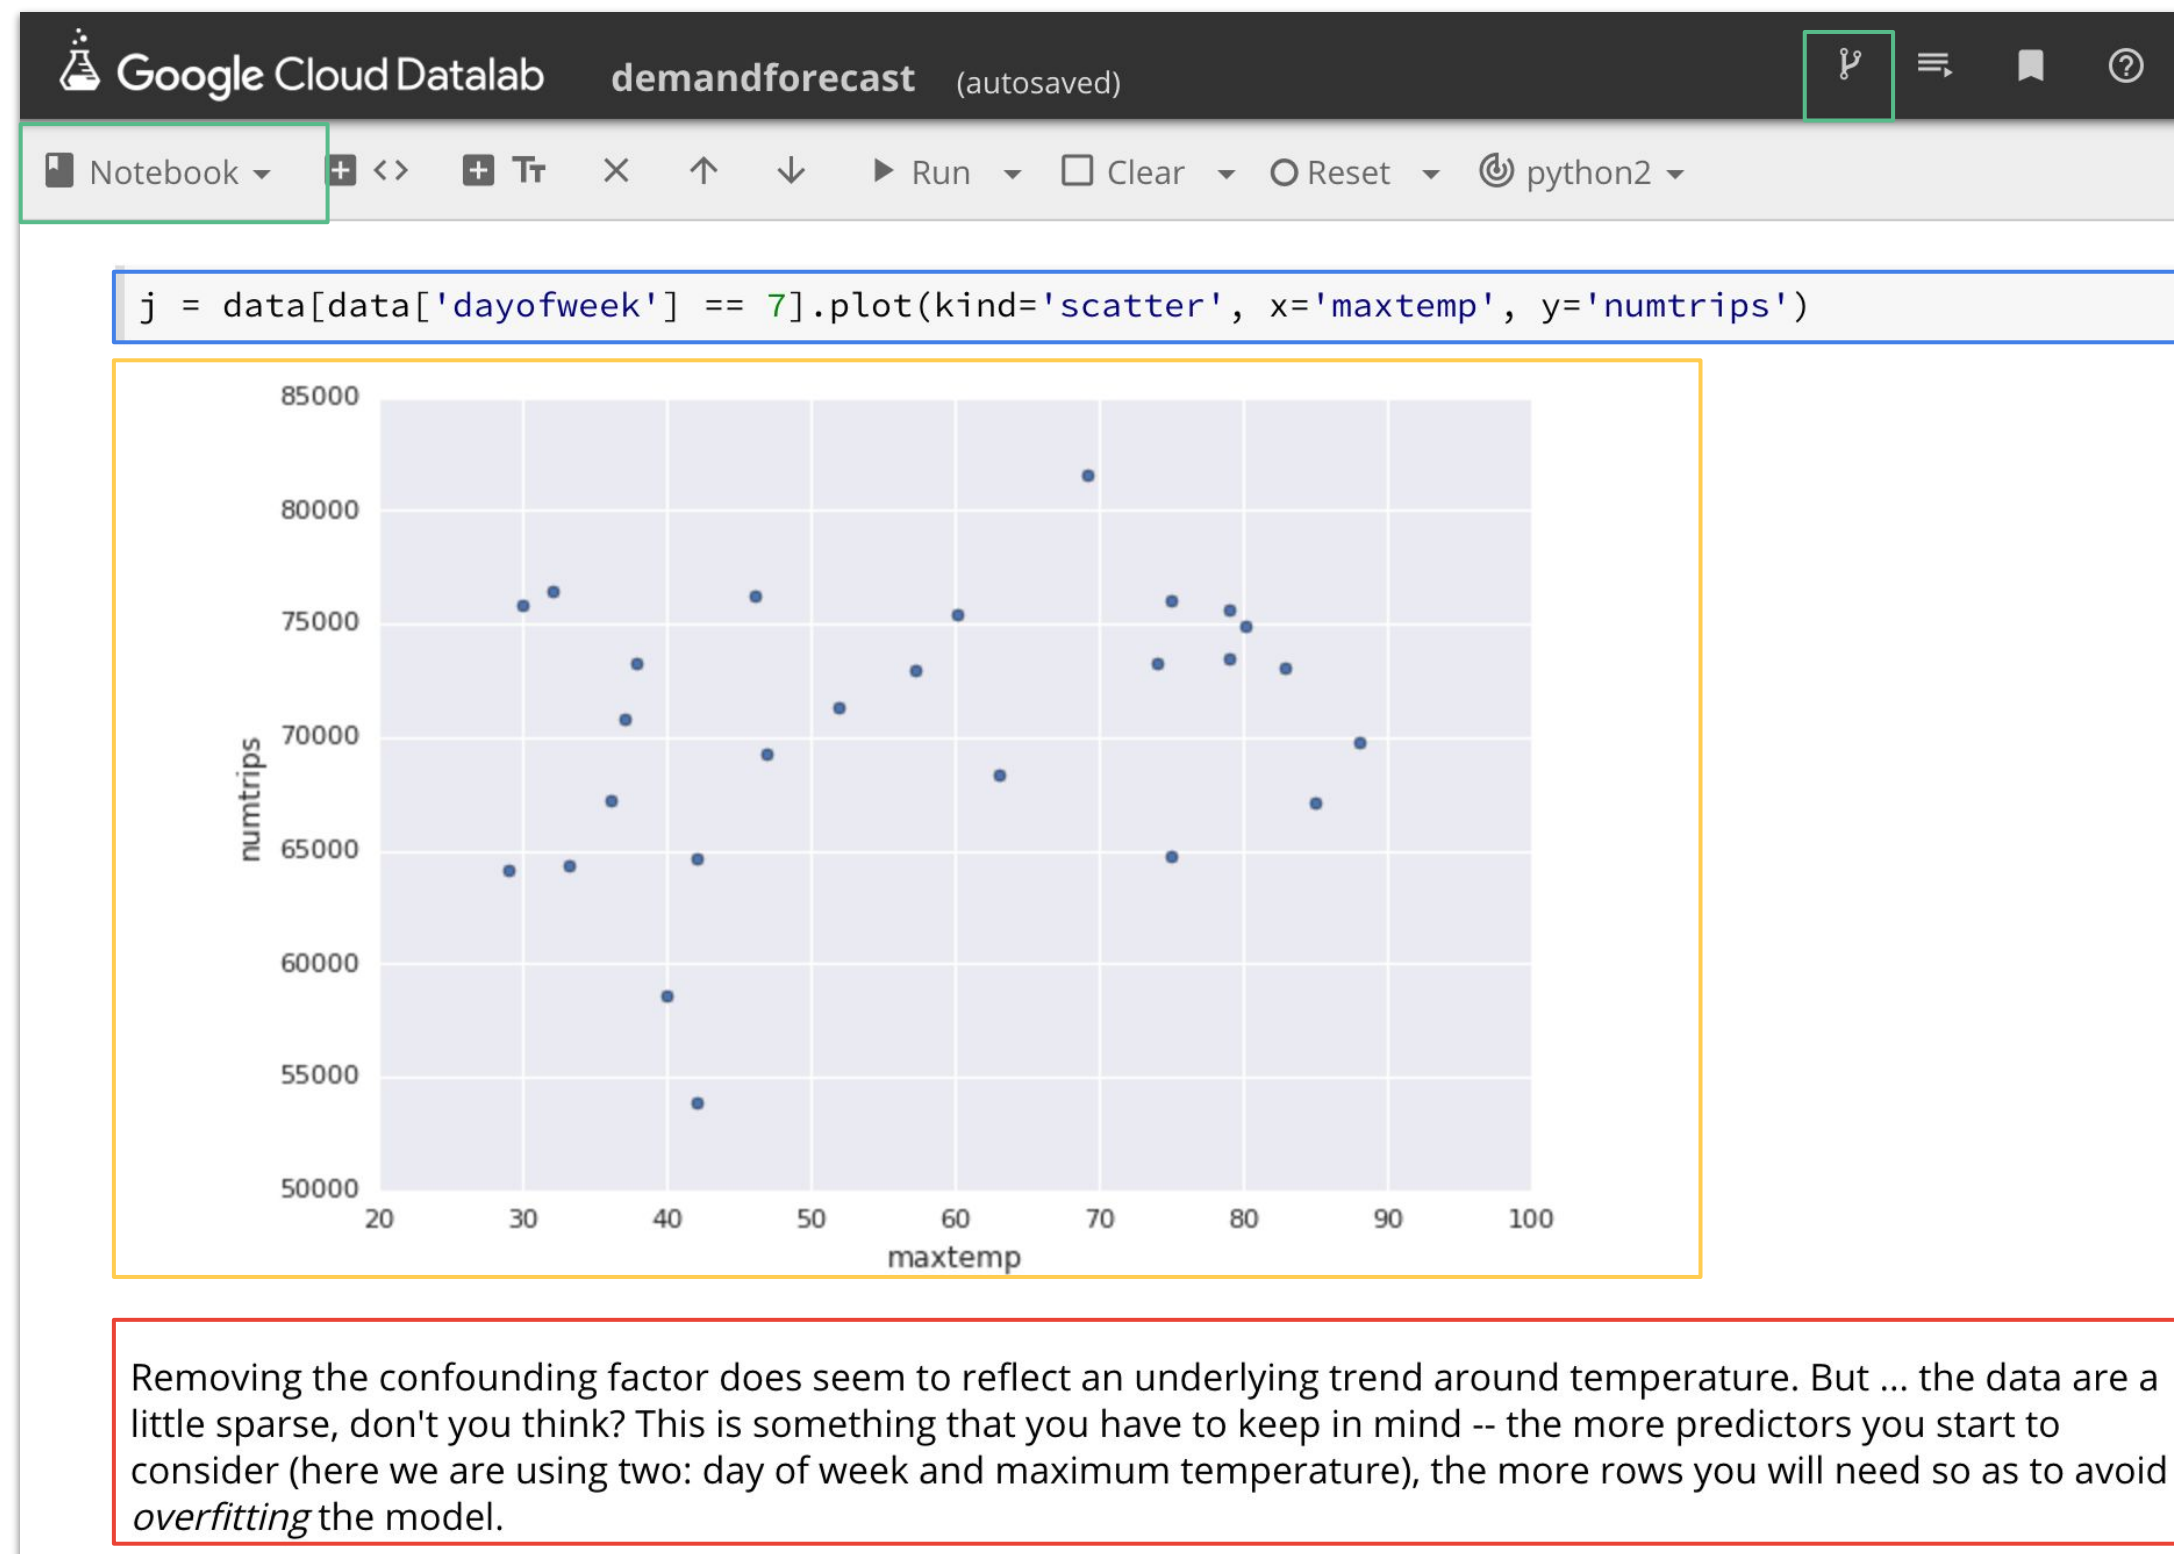
Increasingly, data analysis and machine learning are carried out in self-descriptive, shareable, executable notebooks

Share

Code

Output

Markup



A typical notebook contains code, charts, and explanations

Image Source:
[Git Logo from Wikipedia](#)

Datalab is an open-source notebook built on Jupyter (IPython)

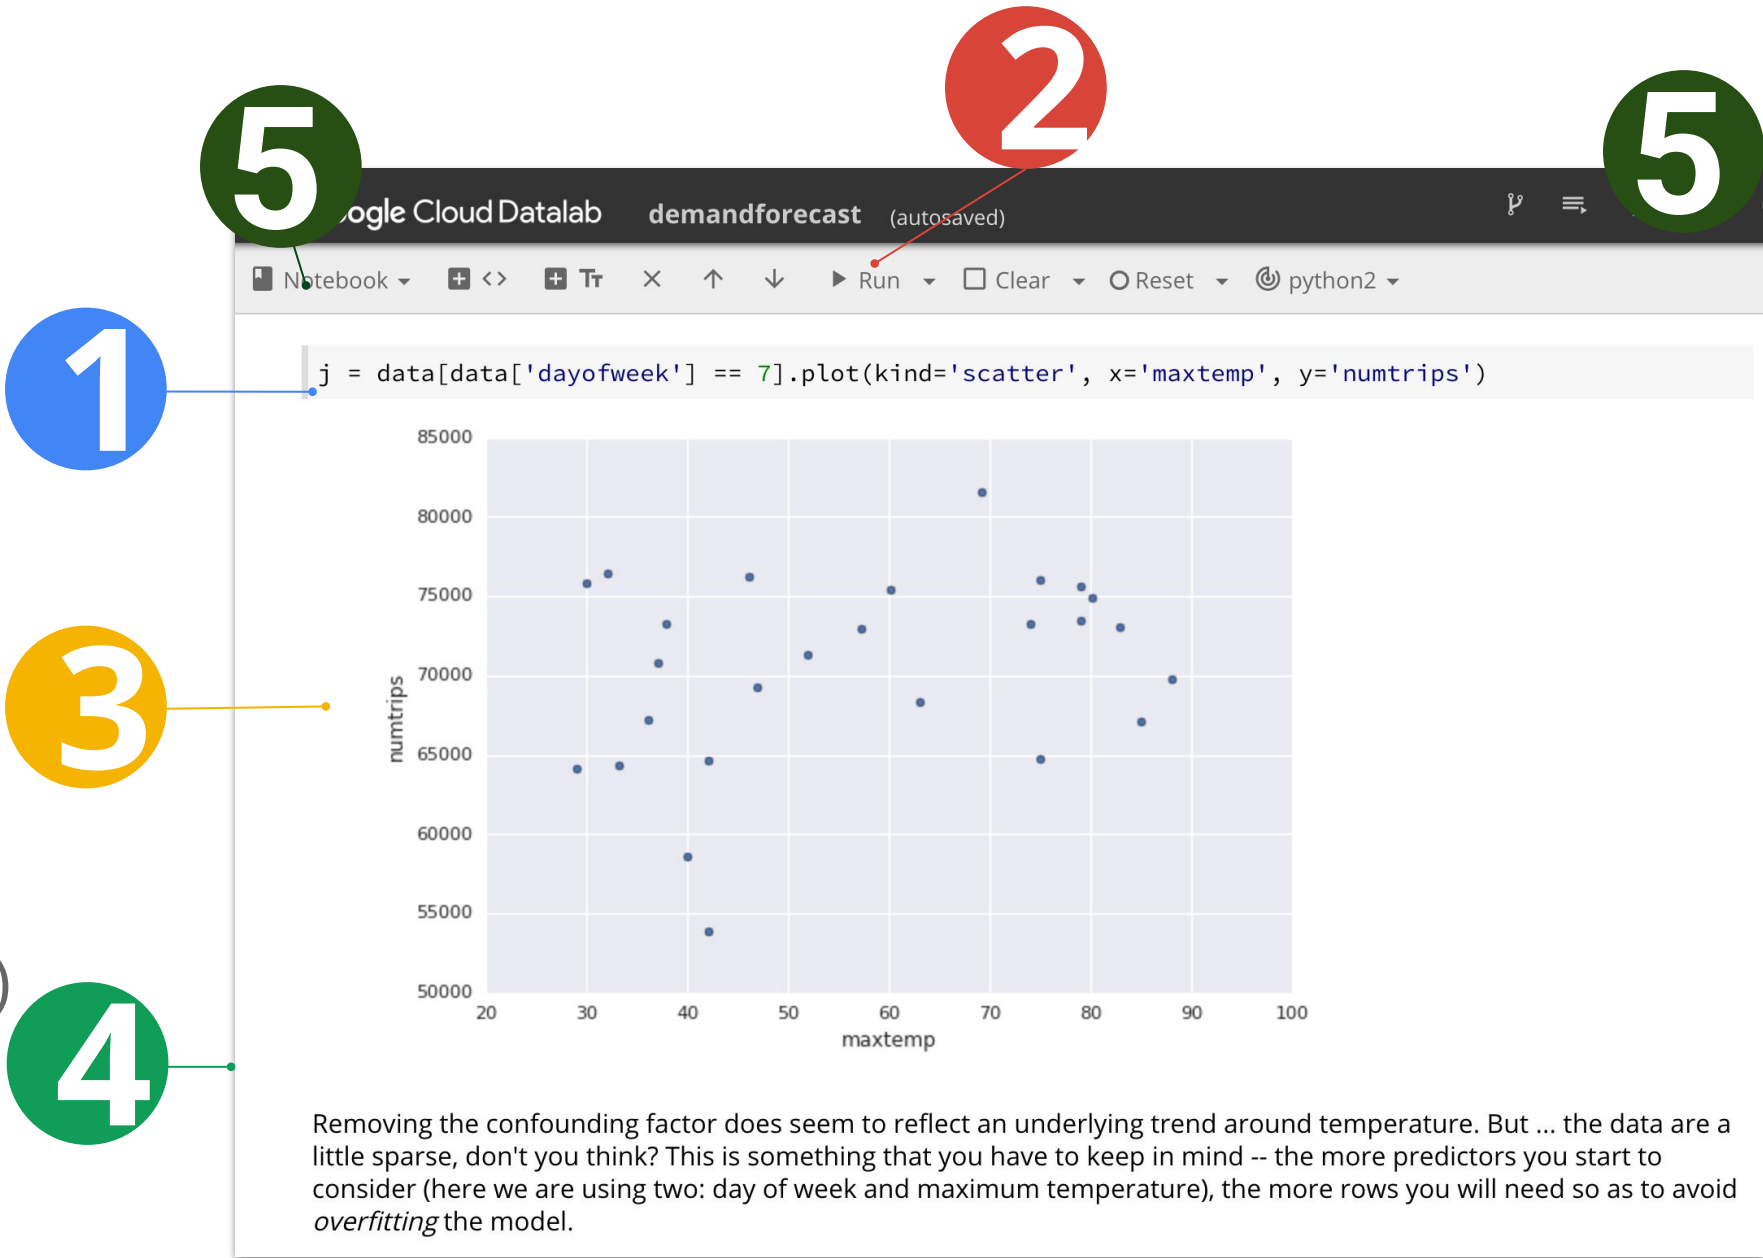
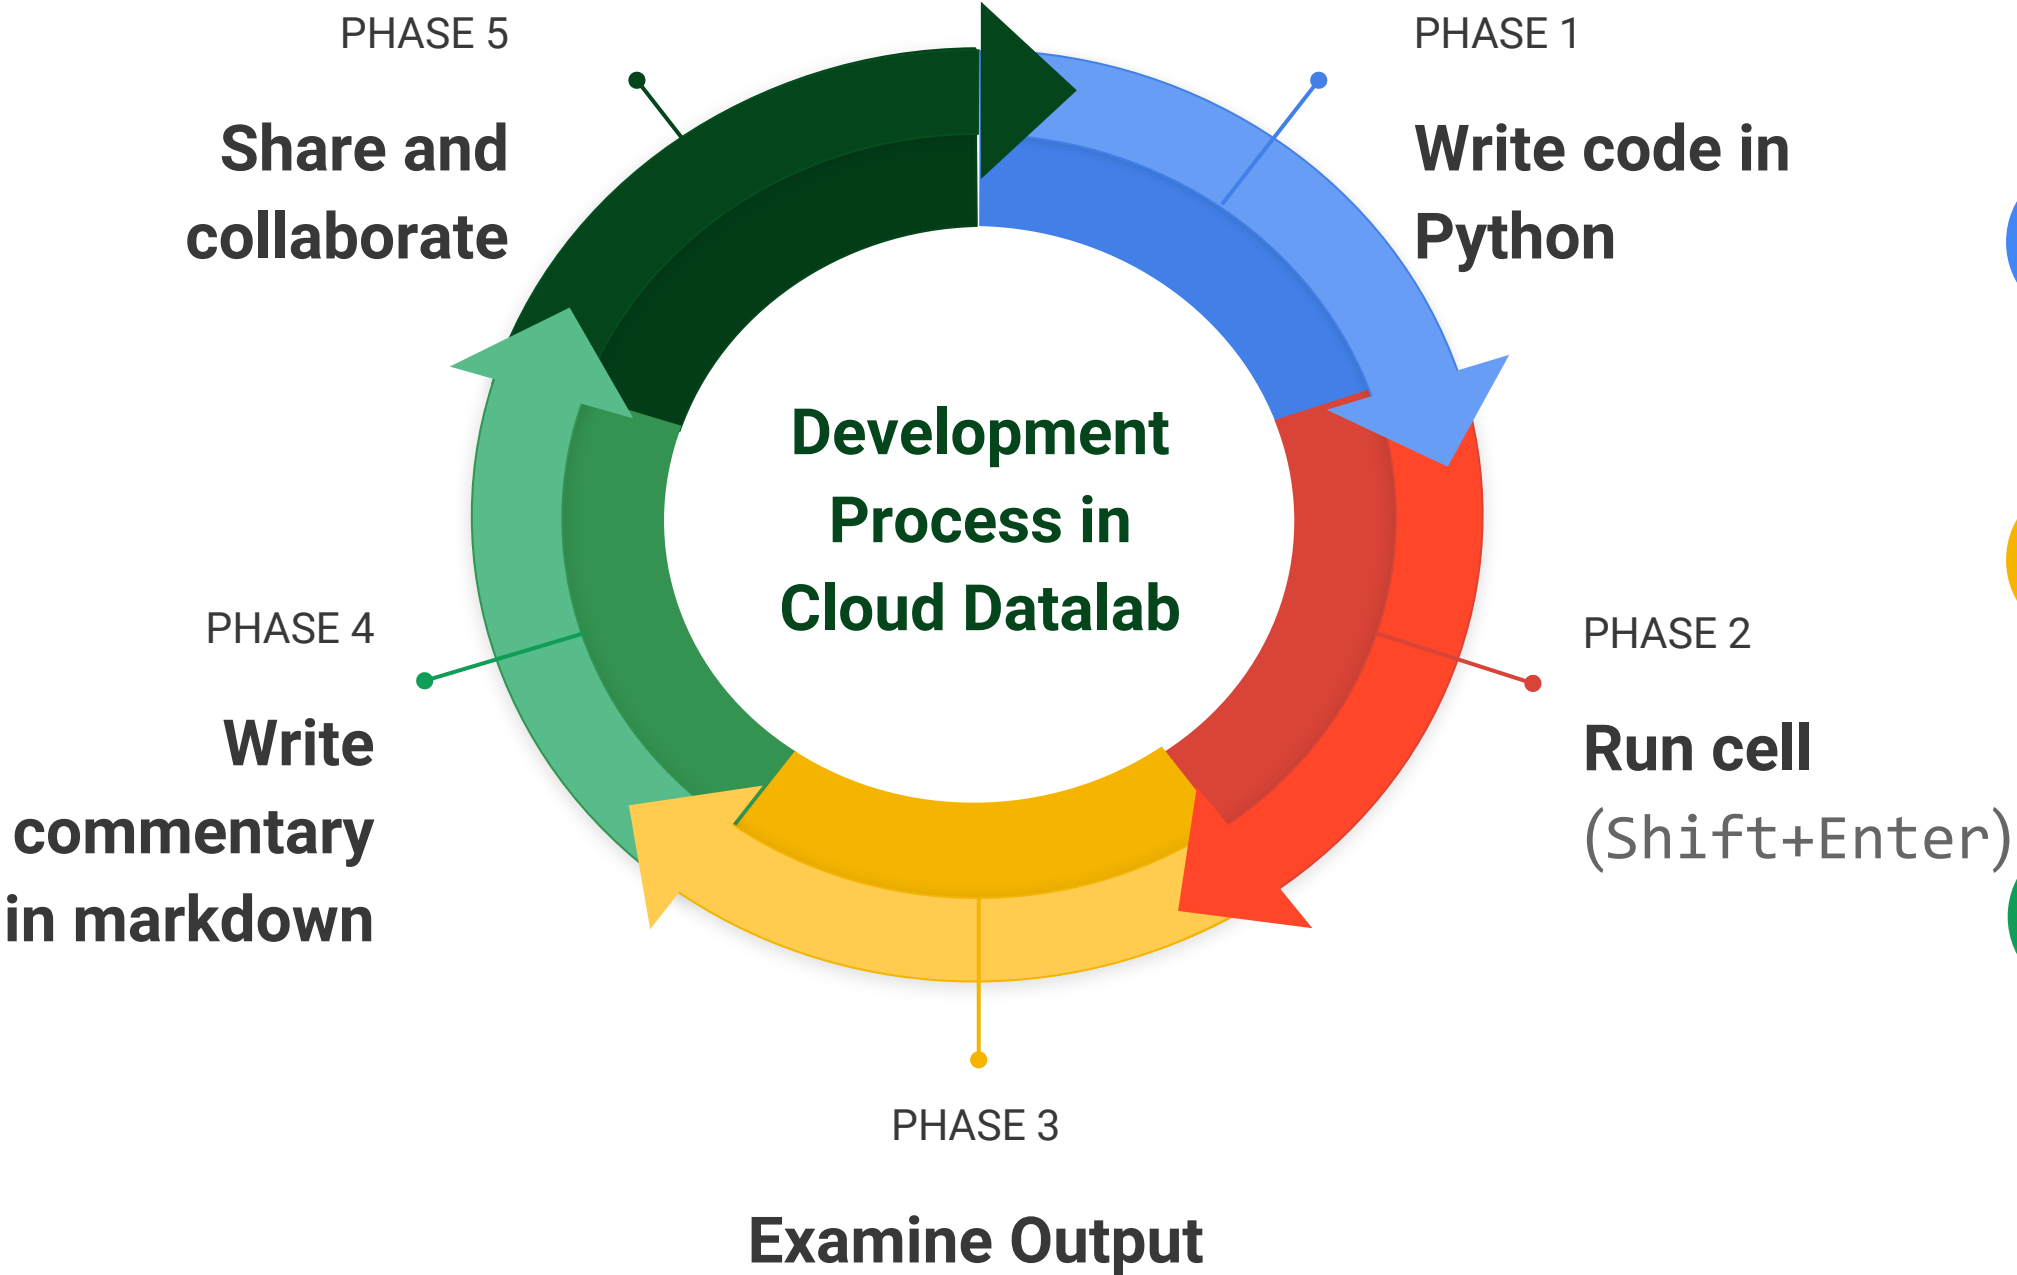
Analyze data in BigQuery,
Compute Engine or Cloud
Storage

Datalab is free—just pay for Google
Cloud resources

Use existing Python
packages



Datalab notebooks are developed in an iterative, collaborative process



Datalab supports BigQuery

```
query = """
SELECT
  weight_pounds,
  is_male,
  mother_age,
  plurality,
  gestation_weeks,
  ABS(FARM_FINGERPRINT(CONCAT(CAST(YEAR AS STRING), CAST(month AS STRING)))
FROM
  publicdata.samples.natality
WHERE year > 2000
"""
```

```
# Call BigQuery and examine in dataframe
import google.datalab.bigquery as bq
df = bq.Query(query + " LIMIT 100").execute().result().to_dataframe()
df.head()
```

	weight_pounds	is_male	mother_age	plurality	gestation_weeks	hashmonth
0	3.562670	True	25	1	30	1403073183891835564
1	3.999185	False	30	1	32	7146494315947640619

Lab: Create ML dataset with BigQuery

Lab 5: Create ML dataset with BigQuery

In this lab, we use BigQuery to create a dataset that we later use to build a taxi demand forecast system using Machine Learning.

- What kinds of things affect taxi demand?
- What are some ways to measure “demand”?



Lab 5: Create ML dataset with BigQuery

In this lab, we use BigQuery to create a dataset that we later use to build a taxi demand forecast system using Machine Learning.

1. Use BigQuery and Datalab to explore and visualize data
2. Build a Pandas dataframe that will be used as the training dataset for machine learning using TensorFlow

Resources

Cloud Datastore

<https://cloud.google.com/datastore/>

Cloud Bigtable

<https://cloud.google.com/bigtable/>

Google BigQuery

<https://cloud.google.com/bigquery/>

Cloud Datalab

<https://cloud.google.com/datalab/>

cloud.google.com